

Part I

Log analysis

1 Introduction

In this assignment you will analyze log files using Apache Spark. We will focus on Linux syslog files and more precisely on the `messages` file located in `/var/log` on most Linux distributions. Such analyses have various applications in system administration, e.g., for security diagnostics.

Lines in `/var/log/messages` are formatted as follows:

```
<timestamp> <host_name> <source>: <message>
```

where `<timestamp>` is a timestamp, `<host_name>` is the host (computer) name, `<source>` is the program that generated the line and `<message>` is the log message. Here is an example:

```
Feb 28 03:30:01 iliad systemd: Started Session 4415 of user achille.
```

You have to write a Spark program to answer the **9 questions** detailed below (**Q1-Q9**). Log files collected on two hosts called `iliad` and `odyssey` are available on Moodle to test your program. During marking, your program will be run on similar but different files to check its correctness. Section 6 summarizes submission instructions and the marking scheme that will be used to evaluate your submission. You are strongly encouraged to review these instructions before starting the assignment.

2 Filtering and Counting

The following queries must be run separately on the log files of each host, i.e., *one answer must be returned for each host*.

Q1 - For each host, print the total number of available log lines.

Expected output on `iliad` and `odyssey`:

```
$ ./log_analyzer -q 1 iliad odyssey
* Q1: line counts
+ iliad: 63854
+ odyssey: 65405
```

Q2 - For each host, print the number of sessions that were started for user `achille`. Session starts are logged in messages containing the following string: `Starting Session <id> of user <user>`, where `<id>` is the session id and `<user>` is the user name.

Expected output on `iliad` and `odyssey`:

```
$ ./log_analyzer -q 2 iliad odyssey
* Q2: sessions of user 'achille'
+ iliad: 5173
+ odyssey: 5228
```

Q3 - For each host, list the unique user names who started a session.

Expected output on iliad and odyssey:

```
$ ./log_analyzer -q 3 iliad odyssey
* Q3: unique user names
+ iliad: ['gaia', 'pollux', 'achille', 'helene', 'hector']
+ odyssey: ['achille', 'hector', 'ares']
```

Q4 - For each host, list the number of sessions started per user.

Expected output on iliad and odyssey:

```
$ ./log_analyzer -q 4 iliad odyssey
* Q4: sessions per user
+ iliad: [('gaia', 2), ('pollux', 38), ('achille', 5173), ('helene', 248), ('hector', 9)]
+ odyssey: [('achille', 5228), ('hector', 2), ('ares', 40)]
```

3 Errors

The following queries must be run separately on the logs of each host, i.e., *one answer must be returned for each host*.

Q5 - For each host, count the error messages, i.e., the lines that contain string “error” (case insensitive match).

Expected output on iliad and odyssey:

```
$ ./log_analyzer -q 5 iliad odyssey
* Q5: number of errors
+ iliad: 2723
+ odyssey: 25805
```

Q6 - For each host, print the 5 most frequent error messages and their counts.

Expected output on iliad and odyssey:

```
$ ./log_analyzer -q 6 iliad odyssey
* Q6: 5 most frequent error messages
+ iliad:
- (889, 'journal: ethtool ioctl error: No such device ')
- (24, 'gnome-session: ** (evince:31187): WARNING **: Error setting file metadata: No such file or directory ')
- (9, 'gnome-session: https://yum.dockerproject.org/repo/main/centos/7/repodata/3849c07e5505140b8134d3cf1bef35c
- (9, "gnome-session: _GDBus.Error:org.gtk.GDBus.UnmappedGError.Quark._imsettings_2derror_2dquark.Code5:_Current
- (8, 'firefox.desktop: Crash Annotation GraphicsCriticalError: |[0][GFX1-]: GLContext is disabled due to a pre
+ odyssey:
- (9229, 'gnome-session: (tracker-miner-fs:30474): Tracker-CRITICAL **: Could not execute sparql: column nie:url
- (4519, 'gnome-session: (tracker-miner-fs:30474): Tracker-CRITICAL **: (Sparql buffer) Error in task 0 of the
- (2776, 'gnome-session: (tracker-miner-fs:30474): Tracker-CRITICAL **: (Sparql buffer) Error in task 2 of the
- (2401, 'gnome-session: (tracker-miner-fs:1259): Tracker-CRITICAL **: Could not execute sparql: column nie:url
- (1697, 'gnome-session: (tracker-miner-fs:1259): Tracker-CRITICAL **: (Sparql buffer) Error in task 0 of the a
```

4 Combining logs

The following queries are run on a combination of both hosts, i.e., *a single answer combining information coming from both hosts must be returned for each question*.

Q7 - List the user names who started a session on both hosts.

Expected output on iliad and odyssey:

```
$ ./log_analyzer -q 7 iliad odyssey
* Q7: users who started a session on both hosts, i.e., on exactly 2 hosts.
+ : ['achille', 'hector']
```

Q8 - List the user names who started a session on exactly one host and list this host.

Expected output on iliad and odyssey:

```
$ ./log_analyzer -q 8 iliad odyssey
* Q8: users who started a session on exactly one host, with host name.
+ : [('gaia', 'iliad'), ('pollux', 'iliad'), ('helene', 'iliad'), ('ares', 'odyssey')]
```

5 Anonymization

Q9 - Anonymize the logs, i.e., replace user names with strings formatted as `user-<i>` where `i` is the index of the user name in the array of original user names sorted alphabetically. For instance, user names `foo` and `bar` must be replaced by `user-1` and `user-0` (respectively), assuming that `foo` and `bar` are the only user names in the log. In addition to writing the anonymized files, your program must print the mapping used for the anonymization and the location where the anonymized files were written.

Expected output on iliad and odyssey:

```
$ ./log_analyzer -q 9 iliad odyssey
+ iliad:
. User name mapping: [('achille', 'user-0'), ('gaia', 'user-1'), ('hector', 'user-2'),\
  ('helene', 'user-3'), ('pollux', 'user-4')]
. Anonymized files: iliad-anonymized-10
+ odyssey:
. User name mapping: [('achille', 'user-0'), ('ares', 'user-1'), ('hector', 'user-2')]
. Anonymized files: odyssey-anonymized-10
```

The corresponding anonymized logs are available on Moodle.